



Degree Project in Computer Science and Engineering

Second cycle, 30 credits

Adaptive Hybrid Retrieval-Augmented Generation

Optimizing the Computational Cost-Performance Trade-off via
Learned Query Routing

DANTE WESSLUND

Adaptive Hybrid Retrieval-Augmented Generation

Optimizing the Computational Cost-Performance Trade-off via Learned Query Routing

DANTE WESSLUND

Master's Programme, Machine Learning, 120 credits

Date: July 5, 2026

Supervisors: Arvid Eriksson, Peng Zhang

Examiner: Pawel Herman

School of Electrical Engineering and Computer Science

Host company: Ericsson

Swedish title: Adaptiv hybrid-RAG

Swedish subtitle: Optimering av beräkningskostnads-prestanda-avvägningen med inlärd frågeroutning

Abstract

Retrieval-Augmented Generation (RAG) is a common approach for grounding large language models in external knowledge. Production systems increasingly combine low-cost vector retrieval with more expensive graph-enhanced retrieval, because graph-enhanced retrieval can help answer questions whose evidence is split across documents but uses substantially more compute than vector retrieval. Applying graph-enhanced retrieval to every query wastes resources on questions that the low-cost path could already handle, while always using vector retrieval risks missing questions that benefit from the richer retrieval mode.

This thesis studies whether the choice between a low-cost vector retrieval mode and a more expensive graph-enhanced retrieval mode can be predicted from the query text alone. The setting is the 2WikiMultihopQA benchmark inside the LightRAG retrieval framework, with LightRAG’s Naive mode as the low-cost path and its Mix mode as the graph-enhanced path. Routing labels are constructed by running both retrieval modes on every benchmark question and using a large language model as a judge to compare each mode’s answer against the gold answer. The resulting binary routing dataset is split jointly by question type and routing label, and is used to train and evaluate two query-only routers: a TF-IDF logistic-regression baseline and a fine-tuned ModernBERT classifier.

On a held-out test split of 193 queries, the ModernBERT router achieves higher point estimates than the TF-IDF baseline across the aggregate routing-label classification metrics (area under the precision-recall curve, balanced accuracy, and minority-class F1), with gains that are stable across three training seeds but small relative to the uncertainty of a test set this size. It also achieves higher F1, precision, and routing-label accuracy than a non-deployable diagnostic heuristic that routes by a coarse question-category label, while using only the query text as input. Treated as offline routing policies, both learned routers route at substantially lower mean cost than always using the Mix mode: the ModernBERT router reduces mean routed execution time by 42 % and estimated mean routed model-token usage by 63 % relative to the always-Mix baseline, while still producing a correct final answer on about 92 % of the queries that always-Mix answers correctly. These figures measure whether the router selects the judge-derived minimum sufficient retrieval mode, not an independently re-judged final answer accuracy of the routed system. Within this experimental setting, the retrieval-mode distinction is partly predictable from the query text alone, and a lightweight encoder

classifier is a promising routing layer for adaptive Hybrid RAG.

Keywords

Retrieval-Augmented Generation, Hybrid RAG, Query routing

Sammanfattning

Retrieval-Augmented Generation (RAG) är en vanlig arkitektur för att förankra stora språkmodeller i externt material. I moderna system kombineras allt oftare vektorbaserad sökning med låg beräkningskostnad med dyrare grafbaserad hämtning, eftersom grafbaserad hämtning kan hjälpa till att besvara frågor vars belägg är fördelade över flera dokument men kräver väsentligt mer beräkningsresurser än vektorbaserad sökning. Att tillämpa den grafbaserade hämtningen på varje fråga slösar resurser på frågor som metoden med lägre kostnad redan klarar, medan att alltid använda enbart vektorsökning riskerar att missa frågor som drar nytta av det mer omfattande hämtningsläget.

Denna avhandling undersöker om valet mellan ett vektorbaserat hämtningsläge med låg kostnad och ett dyrare grafbaserat hämtningsläge kan förutsägas enbart utifrån frågetexten. Experimenten utförs på riktmärket 2WikiMultihopQA i hämtningsramverket LightRAG, där LightRAG:s *Naive*-läge utgör vägen med lägre kostnad och dess *Mix*-läge den grafbaserade vägen. Routingetiketter konstrueras genom att köra båda hämtningslägena på varje fråga och låta en stor språkmodell som domare jämföra respektive svar mot facit. Det resulterande binära routingdatasetet delas upp gemensamt efter frågetyp och routingetikett och används för att träna och utvärdera två routrar som endast ser frågetexten, nämligen en TF-IDF-baserad logistisk regressionsmodell som baslinje och en finjusterad ModernBERT-klassificerare.

På en testmängd om 193 frågor uppnår ModernBERT-routern högre punkttestimat än TF-IDF-baslinjen på de aggregerade routingklassificeringsmått (arean under precision-recall-kurvan, balanserad noggrannhet och F1 på minoritetsklassen), med förbättringar som är stabila över tre träningsfrön men små i förhållande till osäkerheten hos en testmängd av denna storlek. Routern uppnår även högre F1, precision och routingnoggrannhet än en icke-driftsättningsbar diagnostisk heuristik som ruttar utifrån en grov frågekategorietikett, samtidigt som den endast använder frågetexten som indata. Som offline-routingpolicyer ger båda routrarna betydligt lägre medelkostnad än att alltid använda Mix-läget: ModernBERT-routern minskar medellexekveringstiden med 42 % och den uppskattade användningen av modelltoken med 63 % relativt always-Mix-baslinjen, samtidigt som den ger ett korrekt slutligt svar på omkring 92 % av de frågor som always-Mix besvarar korrekt. Dessa mått avser huruvida routern väljer den bedömda minsta tillräckliga hämtningsmetoden, inte en separat ombedömd slutlig svarskorrekthet för hela routingsystemet. I denna experimentella miljö är valet av hämtningsläge delvis förutsägbart enbart utifrån frågetexten, och en

lätteviktig encoder-klassificerare är ett lovande routinglager för adaptiv hybrid-RAG.

Nyckelord

Retrieval-Augmented Generation, hybrid-RAG, frågeroutning

Acknowledgments

I would like to thank my KTH supervisor Arvid Eriksson and my Ericsson supervisor Peng Zhang for their guidance, feedback, and continuous support throughout this thesis project.

Stockholm, July 2026

Dante Wesslund

Contents

1	Introduction	1
1.1	Background	1
1.2	Research problem	2
1.2.1	Problem definition	2
1.2.2	Research question	3
1.3	Purpose	3
1.4	Goals	4
1.5	Research methodology	4
1.6	Delimitations	5
1.7	Structure of the thesis	6
2	Background	7
2.1	Retrieval-augmented generation	7
2.2	Graph-enhanced retrieval	8
2.2.1	Graph-based RAG systems	9
2.2.2	LightRAG <i>Naive</i> and <i>Mix</i> modes	9
2.3	Adaptive retrieval and query routing	10
2.4	Language models and text classifiers	11
2.4.1	Encoder-only transformers and ModernBERT	12
2.4.2	TF-IDF logistic regression	12
2.5	Related work	13
3	Methods	15
3.1	Data	16
3.1.1	2WikiMultihopQA	16
3.1.2	Preprocessing and indexing	16
3.2	Routing-label pipeline	17
3.2.1	Dual-mode retrieval	17
3.2.2	LLM-as-a-Judge labeling	18

3.3	Binary routing dataset	19
3.4	Router models	20
3.4.1	TF-IDF logistic-regression baseline	20
3.4.2	ModernBERT router	21
3.5	Evaluation	22
3.5.1	Held-out classification metrics	22
3.5.2	Offline routing policy	23
4	Implementation	25
4.1	Model training	25
4.1.1	TF-IDF logistic-regression baseline	25
4.1.2	ModernBERT router	26
4.2	Held-out classification experiment	27
4.3	Offline routed-cost experiment	27
4.4	Results	28
4.5	Discussion	34
5	Conclusions and future work	39
5.1	Conclusions	39
5.2	Limitations and future work	41
5.2.1	Token-cost measurement scope	41
5.2.2	Judge labels	41
5.2.3	Final answer accuracy of routed policies	42
5.2.4	Generalizability	42
5.2.5	Routing beyond the query	43
5.2.6	Calibration and abstention	43
5.3	Reflections	43
	References	45
A	Source code	49
B	Experimental configuration	51
B.1	Pipeline scripts	51
B.2	Judge prompt	51

List of Figures

2.1	Stylized illustration of vector retrieval (<i>Naive</i>) versus graph-enhanced retrieval (<i>Mix</i>) on a two-hop question. Vector retrieval matches chunks individually against the query, while the graph-enhanced mode can follow entity-relation edges that connect facts living in different documents.	10
2.2	Query routing between LightRAG’s low-cost <i>Naive</i> path and its more expensive <i>Mix</i> path. The router observes only the query and selects a retrieval mode before retrieval runs.	11
3.1	End-to-end pipeline from 2WikiMultihopQA samples to router evaluation. Each step writes its output to disk so the remaining steps can be re-run without repeating expensive earlier steps.	15
4.1	Precision–recall curves on the held-out test split for the term frequency-inverse document frequency (TF-IDF) baseline and the ModernBERT seed-averaged ensemble. The dashed horizontal line is the test-set base rate of <code>mix_required</code> . . .	30
4.2	Receiver-operating-characteristic curves on the held-out test split.	31
4.3	F1 on <code>mix_required</code> as a function of the decision threshold on the held-out test split. Dots mark the validation-selected threshold for each model: TF-IDF $\tau^* = 0.48$ and the ModernBERT ensemble $\tau^* = 0.44$	31
4.4	Mean routed execution time on the held-out test split as the decision threshold sweeps. The two horizontal dashed lines are the always- <i>Naive</i> and always- <i>Mix</i> reference levels.	32

4.5	Routing-label F1 on the <code>mix_required</code> class versus mean routed execution time on the held-out test split. Each curve traces the operating points obtained by sweeping the decision threshold. Stars mark the validation-selected thresholds. Static <i>always-Naive</i> and <i>always-Mix</i> policies are shown as reference points.	33
4.6	Area under the precision-recall curve (AUPRC) on the <code>mix_required</code> class, broken down by 2WikiMultihopQA question type on the held-out test split. ModernBERT bars are the seed-averaged ensemble. The comparison and inference slices have too few <code>mix_required</code> test rows to give an informative AUPRC and are omitted.	34

List of Tables

3.1	Label-flow from raw benchmark samples to the final binary routing dataset, showing how many queries pass each pipeline stage. The indented rows split the 2000 judged queries into the three judge classes; only <code>naive_enough</code> and <code>mix_required</code> enter the binary dataset, while <code>none_enough</code> (neither mode correct) is excluded.	19
3.2	Routing-dataset composition by split and 2WikiMultihopQA question type (the four types are defined in Section 3.1.1). Each per-type cell gives the number of examples of that type in the split, with the number labeled <code>mix_required</code> in parentheses. Splits come from the joint stratified train/validation/test partition.	20
4.1	Routing-label classification performance on the held-out test split. All accuracy values are routing-label accuracies (agreement with the judge-derived label); <i>Bal. acc.</i> is their class-balanced version and P_{mix} , R_{mix} , FI_{mix} are on the <code>mix_required</code> class at the validation-selected threshold. Bracketed values are 95 % bootstrap intervals over 1 000 test-set resamples; ModernBERT is the seed-averaged ensemble of the three runs, bootstrapped identically to TF-IDF. The constant-score baselines have area under the receiver operating characteristic curve (AUROC) = 0.500 and AUPRC equal to the base rate 0.254; the type-only heuristic is non-deployable as it uses the benchmark’s question-category label.	29

4.2	Offline routed-cost and policy diagnostics on the held-out test split. For each policy the table gives the fraction of queries routed to <i>Mix</i> , the resulting mean per-query time and estimated mean tokens (a mode-weighted average of the per-mode means from the instrumented $N = 15$ subset), the token saving relative to always- <i>Mix</i> , and the absolute false-negative (<i>Under-rout.</i>) and false-positive (<i>Over-rout.</i>) routing-error fractions over all test queries. Under-routing carries the answer-quality risk; over-routing is mainly a cost error. ModernBERT is the seed-averaged ensemble.	29
4.3	Per-question-type breakdown on the held-out test split. Metrics are as in Table 4.1 but computed only on each slice’s rows (threshold metrics at each router’s validation-selected threshold); n and n_{mix} are the slice size and its number of <code>mix_required</code> queries. The comparison and inference slices have too few positives (2 and 1) for stable AUPRC/F1 and are shown for completeness only. ModernBERT is the seed-averaged ensemble.	33
B.1	LightRAG, data, and judge configuration used in this thesis. Each row names one configurable component of the pipeline and gives the exact value used in all reported experiments. <i>Item</i> is the name of the component (model checkpoint, framework version, retrieval parameter, or instrumentation choice); <i>Value</i> is the literal value passed to the corresponding API or library. The same generator model and embedding model are used during indexing and at query time. All Gemini calls use a temperature of 0.0.	52
B.2	Router training configuration used in this thesis. Each row names one router hyperparameter or training-pipeline setting and gives the exact value used in all reported experiments. <i>Item</i> is the name of the hyperparameter or setting; <i>Value</i> is the value passed to scikit-learn (for the TF-IDF baseline) or Hugging Face Transformers (for the ModernBERT router). All three ModernBERT seeds use the values listed here; only the random seed differs across the three runs.	53

B.3	Pipeline stages and the script that implements each stage. <i>Stage</i> names a logical step of the data and modelling pipeline (sample preparation, indexing, dual-mode execution, token instrumentation, judging, dataset construction, router training, evaluation, or figure generation). <i>Script</i> is the path of the Python file in the public code repository (Appendix A) that implements the stage; all paths are relative to the repository root. Running the scripts in the order shown reproduces the data, routers, and figures used in this thesis from the raw 2WikiMultihopQA download.	54
-----	---	----

List of acronyms and abbreviations

AUPRC	area under the precision-recall curve
AUROC	area under the receiver operating characteristic curve
BERT	Bidirectional Encoder Representations from Transformers
GPU	Graphics Processing Unit
LLM	large language model
LLM-as-a-Judge	large language model as a judge
RAG	Retrieval-Augmented Generation
SDG	Sustainable Development Goal
TF-IDF	term frequency-inverse document frequency
UN	United Nations

Chapter 1

Introduction

This chapter introduces the thesis topic and motivates the problem studied in this thesis project. Section 1.1 gives the necessary background in retrieval-augmented question answering and the cost of different retrieval modes. Section 1.2 defines the problem and research question. Section 1.3 states the purpose of the work, Section 1.4 presents the goals, Section 1.5 outlines the research methodology, Section 1.6 defines the delimitations, and Section 1.7 describes the structure of the rest of the thesis.

1.1 Background

Retrieval-Augmented Generation (RAG) has become a common architecture for grounding **large language model (LLM)** outputs in external documents, especially when the relevant information is private, domain-specific, or changes after model pretraining [1]. A typical **RAG** system retrieves text passages from a corpus and conditions the generator on those passages, making the language model less dependent on memorized knowledge.

The simplest retrieval layer is vector similarity search over document chunks. It is efficient and often effective for local factual questions, but it can struggle when a query depends on linking evidence across documents, resolving relationships between entities, or comparing distributed attributes. Multi-hop question answering benchmarks are designed around exactly this regime, in which the answer can only be derived by chaining several supporting facts together. Graph-enhanced **RAG** systems address this by introducing structured representations of entities, relations, or higher-level document organization. Recent systems such as GraphRAG and LightRAG show how graph-based structure can support retrieval beyond vector similarity chunks

[2, 3].

Richer retrieval, however, incurs additional computational cost. Graph construction, graph traversal, larger context assembly, and additional language-model calls all add to the cost of retrieval. At Ericsson, the host company motivating this thesis, where internal technical repositories are large and span many documents, graph-enhanced retrieval may improve answer quality for difficult questions, but applying it to every query can waste computational power on questions that lower-cost vector retrieval already handles [4, 5].

This tension motivates adaptive retrieval. Self-RAG and Adaptive-RAG argue that retrieval behavior should depend on the query rather than follow a fixed policy [4, 6]. The same cost-performance reasoning appears in *model routing*, where a router decides for each query whether a powerful but computationally expensive language model is needed or whether a smaller, less expensive one would be sufficient [5].

The question studied here is narrower than general model routing. Instead of routing between different language models, the project routes between retrieval modes inside a fixed Hybrid **RAG** setting. The low-cost path is LightRAG’s *Naive* mode, which uses only vector retrieval, and the high-cost path is LightRAG’s *Mix* mode, which combines vector retrieval with graph-enhanced retrieval. The routing layer is intentionally lightweight, because if the router itself requires a computationally expensive model, much of the intended efficiency gain is lost.

1.2 Research problem

The problem approached in this thesis is whether the choice between a low-cost vector-based retrieval path and a more expensive graph-enhanced retrieval path in Hybrid **RAG** can be made before retrieval runs, from the query text alone. In LightRAG, this is operationalized as a choice between *Naive* retrieval and *Mix* retrieval. Prior work motivates both graph-based **RAG** and adaptive retrieval [2, 3, 6, 4], but does not quantify how well this retrieval-mode choice can be predicted in advance inside a fixed Hybrid **RAG** framework.

1.2.1 Problem definition

The research problem is to decide, for a given query, whether the low-cost vector retrieval mode is judged sufficient or whether the more expensive graph-enhanced retrieval mode is judged beneficial under the experimental pipeline.

The scientific issue is whether this judged retrieval need is predictable from the query alone. A router that sees only the query text can succeed only if the need is encoded in how the question is phrased. If the need is determined mainly by what the retriever surfaces after retrieval has already run, no query-only router can anticipate it.

The label studied here is operational. A query labeled `mix_required` means that, under the chosen LightRAG configuration and judge prompt, *Mix* was judged correct when *Naive* was not. The label captures whether *Mix* was beneficial under the experimental pipeline, without identifying which component of *Mix* (graph traversal, context assembly, or retrieval breadth) drove the improvement.

1.2.2 Research question

Within a fixed Hybrid **RAG** setting, how well can a lightweight query-only classifier predict a judge-derived routing label for the minimum sufficient retrieval mode, and how does using that classifier as an offline routing policy affect routed cost relative to static *Naive* and *Mix* baselines?

The research question is further split into the following sub-questions:

- How well does a fine-tuned encoder-only classifier predict the judge-derived routing label, compared to a **term frequency-inverse document frequency (TF-IDF)** logistic-regression baseline and a non-deployable coarse question-category heuristic?
- When classifier scores are turned into a thresholded routing policy, how much routed execution time and estimated model-token cost are saved relative to always-*Mix* at the selected routing-label operating point?

1.3 Purpose

The purpose of this work is to reduce the computational cost of Hybrid **RAG** retrieval by making the choice of retrieval mode adaptive, so that the more expensive graph-enhanced mode is reserved for the queries that need it rather than applied to every query. Because a routing layer is only worthwhile if it is itself cheap, a second purpose is to test whether that decision can be made from the query text alone with a lightweight classifier, instead of with an additional expensive model call. In a setting such as Ericsson's, where retrieval runs over large internal corpora, even a partial reduction in how often the expensive mode is invoked translates into a meaningful saving at scale.

1.4 Goals

The main goal is to design and evaluate a learned routing layer that predicts whether a query should use LightRAG’s *Naive* or *Mix* retrieval mode. Four concrete objectives structure the work.

1. Implement a reproducible pipeline for preparing samples from a multi-hop question-answering benchmark, indexing them in LightRAG, and executing both *Naive* and *Mix* retrieval modes.
2. Construct a supervised routing dataset by labeling each query according to the minimum judged sufficient retrieval mode, using a **large language model as a judge (LLM-as-a-Judge)** to compare retrieval outputs against benchmark answers.
3. Train lightweight query-only routers, including a ModernBERT-based classifier and a classical **TF-IDF** logistic-regression baseline, to predict the *Naive*-sufficient versus *Mix*-beneficial routing label.
4. Evaluate the routers both as held-out classifiers, using routing-label metrics such as **area under the precision-recall curve (AUPRC)** and F1 on the *Mix*-beneficial class, and as thresholded offline routing policies measured against always-*Naive*, always-*Mix*, and a non-deployable diagnostic heuristic that routes by a coarse question-category label, on mean routed execution time and estimated mean routed model-token usage.

1.5 Research methodology

This thesis is a quantitative empirical study on a public multi-hop question-answering benchmark. The predictability of retrieval-mode selection from the query text is evaluated by training query-only routers on judge-derived routing labels and obtaining classification metrics on held-out data. These include **AUPRC**, balanced routing-label accuracy, and F1 on the minority class. The same routers are additionally evaluated as thresholded offline routing policies and compared against the static always-*Naive* and always-*Mix* baselines along routed-cost metrics. A non-deployable diagnostic heuristic that routes by a coarse question-category label is included as an additional reference, to measure how much of the routing signal is already captured by such a categorical rule. Bootstrap confidence intervals are used to estimate

the uncertainty of these metrics. Comparisons between the learned routers, the static baselines, and the categorical heuristic are then used to answer the research question.

The reported classification metrics measure agreement with the judge-derived routing label. They are distinct from final-answer accuracy of the selected retrieval path. Routing a *Naive*-sufficient query to *Mix* is counted as a routing-label error because the lower-cost mode was sufficient, although the *Mix* answer for that query may still be correct. This distinction matters when interpreting the static always-*Mix* baseline and the routed-cost results in Chapter 4.

1.6 Delimitations

The scope is deliberately limited to query-only retrieval-mode routing. The router receives the query text as input and does not inspect retrieved documents, intermediate graph state, answer drafts, or model uncertainty after generation. This keeps the routing problem inexpensive and well-defined, but it means the project does not evaluate routers that adapt after retrieval has already begun.

The retrieval setting is LightRAG with two modes, *Naive* and *Mix*. These operationalize the broader distinction between vector retrieval and graph-enhanced retrieval, but results should not be assumed to transfer automatically to every Hybrid RAG framework.

The primary benchmark is 2WikiMultihopQA [7], chosen because it is public, provides supporting evidence, and makes the experiment reproducible without depending on confidential data. Ericsson documentation motivates the industrial problem but is not the core benchmark.

Routing labels come from LLM-as-a-Judge and are treated as judged supervision rather than absolute ground truth. Cases where neither retrieval mode produces a correct answer are excluded from router training.

The offline policy evaluation does not re-judge newly generated answers, because it reuses precomputed *Naive* and *Mix* outputs. Unless stated otherwise, reported accuracy values refer to routing-label accuracy. The two routing-error types are reported separately as the under-routing rate (routing a `mix_required` query to *Naive*) and the over-routing rate (routing a `naive_enough` query to *Mix*), and they have different practical weight, as discussed in Chapter 4.

1.7 Structure of the thesis

Chapter 2 presents the technical and scientific background on RAG, graph-enhanced retrieval, adaptive retrieval, query routing, ModernBERT, and related work. Chapter 3 describes the experimental method, including data preparation, LightRAG execution, label generation, model training, and evaluation design. Chapter 4 describes the implementation, presents the held-out classification and offline routed-cost experiments, reports the results, and discusses what they imply about the predictability of retrieval-mode selection. Chapter 5 concludes with answers to the research question and suggestions for future work.

Chapter 2

Background

This chapter introduces the technical and scientific concepts that the rest of the thesis builds on. Section 2.1 describes RAG, why it is needed for grounding large language model outputs, and how the components fit together. Section 2.2 contrasts vector retrieval with graph-enhanced retrieval, including the two LightRAG retrieval modes used in this thesis. Section 2.3 describes adaptive retrieval and query routing. Section 2.4 introduces the language models and text classifiers used by the router. Section 2.5 positions the thesis in relation to prior work.

2.1 Retrieval-augmented generation

Large language models encode an enormous amount of general knowledge in their parameters, but they have two limitations relevant to industrial question answering. They cannot access information that was not part of their training data, which makes them unable to answer questions over private or company-internal documents, and they have a tendency to hallucinate, generating fluent but factually incorrect answers when asked about content outside their training distribution [1]. RAG addresses both problems by retrieving relevant passages from an external corpus at query time and conditioning the language model on those passages, which both gives the model access to information beyond its training data and lets it ground its answers in citable sources [1, 8].

The original RAG formulation by Lewis et al. [8] treats retrieval and generation jointly as a latent-variable problem. Given a query q , a retriever returns a set $\mathcal{D}_k(q)$ of the k most relevant documents from the corpus, and the generator produces an answer a by marginalizing over which document was

retrieved, as in

$$p(a \mid q) = \sum_{d \in \mathcal{D}_k(q)} p_\theta(a \mid q, d) p_\eta(d \mid q), \quad (2.1)$$

where $p_\eta(d \mid q)$ is the retriever (parameters η) assigning a probability to each candidate document given the query, and $p_\theta(a \mid q, d)$ is the generator (parameters θ) producing the answer given the query and a retrieved document. Many production **RAG** systems use simpler decoding procedures than this marginalization, but Equation 2.1 makes the point that matters here explicit. The documents that the retrieval system surfaces directly determine what the generator can answer.

For industrial question answering the motivation is concrete. At Ericsson, internal technical documents cannot be assumed to appear in any public language-model pretraining corpus, so a **RAG** system is needed to ground answers in company-controlled knowledge. The same is true of many other domain-specific settings, which is why **RAG** has become a standard architecture for grounded question answering over private corpora [1].

Many real questions cannot be answered from a single passage. Multi-hop question answering benchmarks such as HotpotQA, 2WikiMultihopQA, and MuSiQue are constructed around questions whose answers require combining information from multiple supporting documents [7, 9, 10]. For example, a question of the form “*Who is the spouse of the director of film X?*” requires the system first to identify the director of X , then to look up that person, and finally to retrieve their spouse. A retriever that scores each document independently against the query may surface only part of the chain.

2.2 Graph-enhanced retrieval

Most practical **RAG** systems start with vector retrieval. Documents are split into chunks, each chunk is encoded into a dense vector by an embedding model, and at query time the chunks whose vectors are closest in the embedding space to the query’s vector are returned to the generator [8, 11]. This is fast and works well when the answer is contained inside a single passage that resembles the query.

Vector retrieval is weaker when the required evidence is split across multiple documents or depends on relationships between entities. Closeness in the embedding space measures local textual similarity but does not capture the structure that links documents together. Matching the query against individual

chunks can therefore surface passages that are each relevant on their own and still miss the full answer chain.

2.2.1 Graph-based RAG systems

Graph-based **RAG** systems address this limitation by constructing a graph over the corpus before any query is asked. Nodes correspond to entities, chunks, or higher-level summaries, and edges correspond to relations between them. At query time the system can then retrieve evidence by traversing this graph rather than only matching the query against individual chunks.

GraphRAG [2] extracts entities and their relations from each chunk using a language model, builds a knowledge graph over the entire corpus, and groups closely related entities into communities. Each community is summarized into a higher-level description, which can be used to answer queries that span large parts of the corpus. LightRAG [3] follows a similar two-level idea but is designed to be faster and easier to deploy. It extracts an entity-relation graph during indexing, combines it with vector chunks at retrieval time, and exposes several retrieval modes that vary in how heavily they use the graph. LightRAG is the framework used in this thesis because it exposes a low-cost vector-only retrieval mode and a more expensive graph-enhanced mode inside the same system, which makes the two modes directly comparable on identical inputs.

Graph-enhanced retrieval is significantly more computationally expensive than pure vector retrieval. Building the graph requires many language-model calls during indexing for entity and relation extraction, and answering a query involves traversing the graph, assembling a larger context, and often additional language-model calls. Microsoft’s cost discussion for GraphRAG identifies these repeated language-model calls as the dominant cost driver [2, 12]. Graph-enhanced retrieval is therefore best treated as an expensive capability that should be invoked only when the question benefits from it.

2.2.2 LightRAG *Naive* and *Mix* modes

LightRAG exposes several retrieval modes, two of which are used as the routed paths in this thesis [3]. The *Naive* mode retrieves the top- k text chunks for the query by vector similarity and passes them directly to the generator, ignoring the graph. The *Mix* mode is hybrid: it retrieves vector chunks as in *Naive*, additionally retrieves entities and relations from the entity-relation graph, and assembles both kinds of evidence into the generator’s context. The two modes share the same underlying corpus and generator but differ in how much graph

structure they use, which makes them a natural pair of routed retrieval paths.

Figure 2.1 illustrates the difference with a stylized multi-hop example. The *Naive* mode retrieves chunks similar to the query and can fail when no single chunk contains both required hops. The *Mix* mode follows entity-relation edges in the graph and can connect the two hops even when they live in separate documents.




Figure 2.1: Stylized illustration of vector retrieval (*Naive*) versus graph-enhanced retrieval (*Mix*) on a two-hop question. Vector retrieval matches chunks individually against the query, while the graph-enhanced mode can follow entity-relation edges that connect facts living in different documents.

2.3 Adaptive retrieval and query routing

Adaptive retrieval is motivated by the fact that different queries benefit from different retrieval strategies. Self-RAG [6] fine-tunes the generator to emit reflection tokens that decide whether retrieval is necessary, what to retrieve, and whether the generated answer is supported by the evidence, which lets the system skip retrieval entirely on questions it is confident about. Adaptive-RAG [4] takes a different approach: a small classifier estimates the complexity of a question and the system uses that estimate to select between progressively more expensive retrieval-augmented reasoning strategies, ranging from no retrieval through single-hop retrieval to iterative retrieval. Both works support the premise that retrieval behavior should depend on the query, and Adaptive-RAG in particular shows that a small upstream classifier can drive the strategy choice.

Model routing frames the same idea more generally as a cost–quality decision [5]. In a multi-model setting, the system has access to several language models with different cost and accuracy profiles, and a router decides on a per-question basis whether to use a powerful but expensive model or a smaller, less expensive one. RouterBench [5] is a benchmark and evaluation suite for this problem and reports that even simple routers can recover a large fraction of the strongest model’s accuracy while spending only a fraction of its cost. Although RouterBench concerns routing between language models rather than between retrieval modes, the abstraction transfers directly. A router observes an input and chooses among execution paths with different cost and expected performance before the expensive computation happens.

This thesis applies that routing abstraction inside the retrieval layer, selecting between two retrieval modes within one Hybrid **RAG** framework. Figure 2.2 illustrates the routing decision. Adaptive-RAG and similar work likewise make their routing decision before retrieval runs, on the query alone [4, 5].




Figure 2.2: Query routing between LightRAG’s low-cost *Naive* path and its more expensive *Mix* path. The router observes only the query and selects a retrieval mode before retrieval runs.

2.4 Language models and text classifiers

The routing layer is implemented as a text classifier over the query. This section describes the two classifier families compared in this thesis, a transformer-based encoder model (ModernBERT) and a classical **TF-IDF** logistic-regression baseline.

2.4.1 Encoder-only transformers and ModernBERT

The transformer architecture [13] introduced self-attention as a general mechanism for contextualized text representations. Encoder-only models such as **Bidirectional Encoder Representations from Transformers (BERT)** are trained on masked language modeling and produce a contextualized vector per input token, which can be pooled or fed to a small classification head to produce a class label [14]. Decoder-only models such as the GPT family are instead trained to predict the next token and are more naturally used as generative language models.

Both families can in principle be used for classification, but encoder-only models remain attractive when the output is a class label and inference efficiency matters. A recent comparison by Benayas et al. [15] finds that encoder-only models match or exceed much larger decoder-only models on intent classification and sentiment analysis while being orders of magnitude smaller, which is consistent with the design goal of a low-cost routing layer.

ModernBERT [16] is a recent encoder model that updates the **BERT** architecture with modern training data, longer context, and several efficiency improvements. The authors report that ModernBERT-base matches or exceeds the classification quality of much larger models while remaining light enough to fine-tune on a single **Graphics Processing Unit (GPU)**, which makes it a natural choice for a query-only router that has to remain inexpensive relative to the retrieval modes it routes between.

2.4.2 TF-IDF logistic regression

The classical baseline used in this thesis is a logistic regression over **TF-IDF** features [17]. Each query is represented by a sparse vector whose entries are term frequencies weighted by the inverse document frequency of the term in the training corpus, which downweights very common words. A logistic regression then maps this sparse feature vector to a probability that the query is *Mix*-beneficial under the judged setup. The baseline matters because the distinction between *Naive*-sufficient and *Mix*-beneficial queries may be partly lexical. Multi-hop questions often have characteristic surface patterns. Nested relations like “*the spouse of the director of X*” signal a chained lookup, and comparisons between two entities signal that evidence must come from multiple sources. A **TF-IDF** classifier can pick up such patterns from word frequencies alone, so it sets a meaningful floor for how much a pretrained transformer encoder needs to add.

2.5 Related work

The most closely related prior work spans graph-based **RAG** systems, adaptive-retrieval work that uses an upstream classifier, and general model routing.

Edge et al. [2] introduce GraphRAG, and Guo et al. [3] introduce LightRAG. Both systems demonstrate that graph-enhanced retrieval can add structure beyond local vector similarity and improve answer quality on questions that require linking facts across documents. They also document the additional cost of doing so. Neither work addresses the follow-up question of how the system should decide which retrieval mode to use for a given query. This thesis treats that selection as a supervised prediction problem.

Jeong et al. [4] present Adaptive-RAG, the closest precedent on the methodological side. They train a small classifier on the question to estimate complexity and use that estimate to select between progressively more expensive retrieval-augmented reasoning strategies, ranging from no retrieval to iterative retrieval. The classifier is trained on labels derived from how well different strategies answer the question on a training set, which is conceptually similar to the judge-derived labels used here. Adaptive-RAG and this thesis differ in two respects. Adaptive-RAG selects across these reasoning strategies at the language-model level, whereas this thesis selects between LightRAG’s *Naive* and *Mix* retrieval modes inside one fixed framework. The classifier here is also intentionally restricted to a lightweight encoder so the routing cost remains negligible. Asai et al. [6] achieve a similar adaptive behavior with Self-RAG by training the generator itself to emit reflection tokens, which is a more invasive change and does not give an explicit pre-retrieval routing decision.

Hu et al. [5] formalize routing between language models as a cost and quality trade-off and provide RouterBench as a benchmark for it. Although RouterBench routes between language models, the framing carries over directly to retrieval-mode routing. The trade-off plots used in Chapter 4 follow the same logic. Each operating point of the learned router can be compared to static baselines, and the question of interest is whether the router achieves useful routing-label performance at substantially lower routed cost than the always-expensive policy.

This literature does not contain a focused empirical study of query-only routing between a low-cost vector retrieval mode and a more expensive graph-enhanced retrieval mode inside a single Hybrid **RAG** framework, treated as a supervised classification problem with judge-derived labels. This thesis

addresses that gap, using LightRAG's *Naive* and *Mix* modes as the concrete low-cost and graph-enhanced paths.

Chapter 3

Methods

This chapter describes the data, the routing pipeline, the router models, and the evaluation procedure used to study query-only retrieval-mode routing. Section 3.1 describes the benchmark data and how it is preprocessed and indexed. Section 3.2 describes the routing-label pipeline that turns raw benchmark questions into a supervised routing dataset. Section 3.3 describes the resulting binary routing dataset and the held-out splits (the validation and test splits, neither of which is used to fit any model parameters) used for evaluation. Section 3.4 describes the two router models, a ModernBERT-based classifier and a **TF-IDF** logistic-regression baseline. Section 3.5 describes the held-out classification metrics and the offline routing-policy evaluation.

Figure 3.1 gives a visual overview of the full pipeline; the sections that follow describe each step in turn.




Figure 3.1: End-to-end pipeline from 2WikiMultihopQA samples to router evaluation. Each step writes its output to disk so the remaining steps can be re-run without repeating expensive earlier steps.

3.1 Data

This section describes the 2WikiMultihopQA benchmark (Section 3.1.1) and the preprocessing and indexing steps that prepare it for the routing-label pipeline (Section 3.1.2).

3.1.1 2WikiMultihopQA

The benchmark used in this thesis is 2WikiMultihopQA [7], a public multi-hop question answering dataset constructed from Wikipedia. Each sample contains a question, a gold answer (the reference answer supplied by the benchmark), a set of context documents drawn from Wikipedia, and a set of supporting facts that mark which sentences in the context are required to derive the answer. Questions are annotated with one of four types that describe the reasoning structure. These types are `comparison`, `compositional`, `bridge_comparison`, and `inference`.

2WikiMultihopQA is chosen because it is public, which makes the experiment reproducible, and because it is multi-hop by construction, so the answer cannot in general be read off a single passage. The multi-hop property is the setting in which graph-enhanced retrieval is expected to differ most from vector retrieval. The supporting-fact annotations also provide traceable evidence metadata, although this thesis uses them only for dataset traceability and does not perform a full supporting-fact analysis. Other multi-hop benchmarks such as HotpotQA and MuSiQue [9, 10] share the same motivation. Evaluating all three benchmarks jointly would have been preferable, but indexing and dual-mode execution are costly, so in the interest of compute and time a single benchmark was used. 2WikiMultihopQA was preferred because, in preliminary testing across the three benchmarks under LightRAG’s dual-mode execution, it produced the most balanced `mix_required` versus `naive_enough` distribution, whereas the other two gave more skewed distributions that would have made supervised routing harder to evaluate; extending the study to them is left to future work.

3.1.2 Preprocessing and indexing

Each sample is converted into a structured record containing a stable sample identifier, the question, the gold answer, the question type, the context documents, and the supporting facts. The mapping between sample identifier and these fields is preserved throughout the pipeline so that every later artifact

(retrieval output, judge label, router prediction) can be traced back to the original sample.

Context documents from these structured records are then inserted into LightRAG [3], which during indexing builds both a vector index over chunked text and an entity-relation graph derived from the same chunks. Once the index has been built, it is treated as a fixed part of the experimental setup. Stochasticity in graph construction (entity extraction, relation extraction, summarization) is therefore fixed in the index and does not introduce additional variance during the dual-mode retrieval experiments.

3.2 Routing-label pipeline

Figure 3.1 shows the end-to-end pipeline. Starting from 2WikiMultihopQA samples, the corpus is indexed in LightRAG, every sample is executed with both *Naive* and *Mix* retrieval, an **LLM-as-a-Judge** assigns a routing class to each query, and the resulting binary labels are used to train and evaluate query-only routers.

3.2.1 Dual-mode retrieval

For every benchmark question q_i , the indexed corpus is queried with both LightRAG retrieval modes, giving

$$r_i^N = R_N(q_i), \quad r_i^M = R_M(q_i), \quad (3.1)$$

where R_N denotes LightRAG *Naive* (the low-cost vector path) and R_M denotes LightRAG *Mix* (the graph-enhanced path). The recorded artifacts for each query are the two generated answers r_i^N and r_i^M , the measured execution time t_i^N and t_i^M for each mode, and the number of context documents and supporting facts associated with the sample.

Two cost dimensions are reported. Execution time is logged per query for every query in the dataset. Model-token usage is additionally measured on an instrumented subset of $N = 15$ queries spanning all four 2WikiMultihopQA question types, with token counts taken from Gemini API `usage_metadata` when available and otherwise estimated with the configured tokenizer. The instrumented subset is used to estimate the per-mode mean model-token consumption, denoted $\bar{u}^N$ for the *Naive* retrieval mode and $\bar{u}^M$ for the *Mix* retrieval mode. These per-mode means are then applied to the held-out test split for the routed-cost analysis in Section 3.5.2.

The routed token figures are therefore estimated from per-mode token means and are not measured per query for every test example. Section 5.2.1 discusses the implications of this small instrumented subset. Reporting both cost dimensions matters because execution time is wall-clock time that includes the network round-trips of the API-based LightRAG and judge calls, not just model inference, so it would shrink with a locally hosted model; model-token usage, by contrast, corresponds more directly to per-query operational cost in production deployments and is insensitive to network conditions. Both LightRAG retrieval modes generate their answers with the same instruction-tuned Gemini model, and the **LLM-as-a-Judge** uses that same model; all model calls run at temperature 0 for reproducibility, and retrieval uses a Gemini embedding model over the chunked corpus. Using one capable but inexpensive model for both generation and judging keeps the pipeline simple and keeps the judge’s reading of an answer consistent with how that answer was produced. The exact generator, judge, and embedding model checkpoints, the LightRAG version, retrieval parameters, token logging method, and judge prompt are listed in Appendix B.

3.2.2 LLM-as-a-Judge labeling

The 2WikiMultihopQA benchmark provides gold answers, but the routing labels needed here are specific to the two retrieval modes being compared and must therefore be constructed. They are constructed by running an **LLM-as-a-Judge** [18] that compares each query’s two generated answers against the benchmark answer. The judge assigns one of three classes to each query. The class `naive_enough` is assigned when r_i^N matches the gold answer, so the low-cost mode was judged sufficient. The class `mix_required` is assigned when r_i^N does not match the gold answer but r_i^M does, so *Mix* was judged beneficial relative to *Naive* under the experimental pipeline. The class `none_enough` is assigned when neither mode produced an answer that matches the gold answer.

The binary routing label $y_i \in \{0, 1\}$ is then defined as

$$y_i = \begin{cases} 0, & \text{judge class is } \text{naive_enough}, \\ 1, & \text{judge class is } \text{mix_required}, \end{cases} \quad (3.2)$$

with `none_enough` cases excluded from the binary router training set. For `none_enough` queries neither retrieval mode succeeds, so the failure is one of answer quality and lies outside the scope of routing.

The label `mix_required` is operational: it records that *Mix* was judged beneficial relative to *Naive* under this LightRAG setup, without identifying which component of *Mix* (graph traversal, context assembly, or retrieval breadth) drove the improvement.

The **LLM-as-a-Judge** is treated as judged supervision and is not used as ground truth. Three design choices reduce the risk of judge error. The judge sees the gold answer in addition to the two candidate answers, which turns the task into a constrained comparison against a reference instead of open-ended quality scoring and makes the judgment easier and more reliable [18]. The same judge model and prompt are used across the entire dataset, which keeps the labeling protocol internally consistent.

3.3 Binary routing dataset

Before the binary dataset is analyzed by split, Table 3.1 summarizes how the raw dual-mode runs are filtered into the final binary routing dataset.

Table 3.1: Label-flow from raw benchmark samples to the final binary routing dataset, showing how many queries pass each pipeline stage. The indented rows split the 2 000 judged queries into the three judge classes; only `naive_enough` and `mix_required` enter the binary dataset, while `none_enough` (neither mode correct) is excluded.

Stage	Count
Queries successfully run with both <i>Naive</i> and <i>Mix</i>	2 000
Judge produced usable label	2 000
<code>naive_enough</code>	959
<code>mix_required</code> (<i>Mix</i> -beneficial)	327
<code>none_enough</code> (excluded from binary dataset)	714
Final binary routing dataset	1 286

After the dual-mode execution and judge labeling, the binary routing dataset contains 1 286 judged queries with a routing class of either `naive_enough` or `mix_required`. Of these, 959 (74.6 %) are labeled `naive_enough` and 327 (25.4 %) are labeled `mix_required`. The class distribution is therefore imbalanced, which is taken into account during training and during reporting (Section 3.5).

The dataset is split into 900 training examples, 193 validation examples, and 193 test examples (approximately 70/15/15), jointly stratified by question

type and routing label so that both the class proportions and the per-question-type proportions appear in all three splits in the same proportion as in the overall dataset. The validation split is held out from training and is used for early stopping and threshold selection. The test split is held out from both training and threshold selection and is used for final reported metrics.

Table 3.2 summarizes the routing dataset by split and 2WikiMultihopQA question type. The class proportions are preserved across splits by construction, but the distribution of `mix_required` labels is uneven across question types. `Compositional` questions are the most likely to be judged *Mix*-beneficial, while `comparison` questions are usually answered well enough by the low-cost mode.

Table 3.2: Routing-dataset composition by split and 2WikiMultihopQA question type (the four types are defined in Section 3.1.1). Each per-type cell gives the number of examples of that type in the split, with the number labeled `mix_required` in parentheses. Splits come from the joint stratified train/validation/test partition.

Split	Total	<code>comparison</code>	<code>compositional</code>	<code>bridge comp.</code>	<code>inference</code>
Train	900	415 (8 mix)	305 (178 mix)	155 (37 mix)	25 (6 mix)
Validation	193	89 (2 mix)	66 (38 mix)	33 (8 mix)	5 (1 mix)
Test	193	90 (2 mix)	65 (38 mix)	33 (8 mix)	5 (1 mix)
Overall	1 286	594 (12 mix)	436 (254 mix)	221 (53 mix)	35 (8 mix)

3.4 Router models

Both router models take only the question text q_i as input and output a score $s_\phi(q_i) \in [0, 1]$, where ϕ denotes the model parameters. The score is interpreted as $P_\phi(y_i = 1 \mid q_i)$, the model’s estimated probability that the binary routing label y_i is 1, i.e., that *Mix* is judged beneficial for query q_i under the experimental pipeline.

3.4.1 TF-IDF logistic-regression baseline

The baseline router is a logistic regression trained on **TF-IDF** features of the question. The pipeline first builds a vocabulary of V terms from the training questions, where each term is a unigram or bigram after lowercasing and basic English stop-word removal. Each question is then represented as a sparse vector $\mathbf{x}_i \in \mathbb{R}^V$ whose entries are the term frequency of each vocabulary

term multiplied by the inverse document frequency of that term in the training corpus. The logistic-regression head produces the router score as

$$s_\phi(q_i) = \sigma(\mathbf{w}^\top \mathbf{x}_i + b), \quad \sigma(z) = \frac{1}{1 + e^{-z}}, \quad (3.3)$$

where the weight vector $\mathbf{w} \in R^V$ and bias $b \in R$ are learned by minimizing the L_2 -regularized binary cross-entropy loss with balanced class weights, so that the minority `mix_required` class is not ignored.

The baseline is intentionally simple. It has no pretrained semantic representations and can only exploit lexical patterns in the question. Its role is to indicate how much of the routing signal is captured by surface lexical cues, which calibrates how much improvement a transformer encoder can plausibly add.

3.4.2 ModernBERT router

The transformer router is a ModernBERT-base [16] encoder with a binary classification head. The question is tokenized and encoded by ModernBERT into a sequence of contextualized hidden states. The pooled representation $h_i \in R^d$, where d is the ModernBERT hidden dimension, is then mapped to two logits by the classification head,

$$z_i = Wh_i + b, \quad W \in R^{2 \times d}, \quad b \in R^2, \quad (3.4)$$

where logit $z_{i,0}$ corresponds to `naive_enough` and $z_{i,1}$ to `mix_required`. The router score is the softmax probability of the `mix_required` class,

$$s_\phi(q_i) = \frac{e^{z_{i,1}}}{e^{z_{i,0}} + e^{z_{i,1}}}. \quad (3.5)$$

Equivalently, the same probability can be written using a single logit and the sigmoid function, $s_\phi(q_i) = \sigma(z_{i,1} - z_{i,0})$. The two-logit parameterization used here matches the default Hugging Face classification-head format for binary classification with `num_labels=2`.

The model is fine-tuned by minimizing class-balanced cross-entropy on the n training examples,

$$\mathcal{L}(\phi) = -\frac{1}{n} \sum_{i=1}^n \alpha_{y_i} \log P_\phi(y_i | q_i), \quad (3.6)$$

where $\alpha_c = n/(2n_c)$ for class $c \in \{0, 1\}$, with n_c the number of training examples of class c . This weighting balances the contribution of the two classes so that the minority `mix_required` class is not down-weighted by the imbalanced training distribution. Optimization uses AdamW [19] with a linear warm-up and linear decay schedule. The best checkpoint is selected on the validation split using the F1 score on the `mix_required` class, and that checkpoint is then used for all reported test metrics. Hyperparameters are listed in Chapter 4.

3.5 Evaluation

Evaluation separates routing-label prediction from routed-cost analysis. Held-out classification metrics measure agreement with the judge-derived routing label. Routing a `naive_enough` query to *Mix* counts as a routing-label error because the lower-cost mode was already sufficient, while routing a `mix_required` query to *Naive* counts both as a routing-label error and as an answer-quality risk under the label definition.

Offline policy evaluation measures what happens when the router’s predictions are used to select between the precomputed *Naive* and *Mix* retrieval outputs. Both evaluations use the held-out test split. The validation split is used only for early stopping and for selecting an operating threshold.

3.5.1 Held-out classification metrics

Because the binary class distribution is imbalanced, the primary threshold-independent metric is the **AUPRC** on the `mix_required` class. **AUPRC** measures how well the router ranks `mix_required` queries above `naive_enough` queries and is more informative than **area under the receiver operating characteristic curve (AUROC)** when the positive class is in the minority [20]. **AUROC** is also reported for comparison with prior work.

For thresholded evaluation, the router score $s_\phi(q_i)$ is converted to a class prediction by applying a decision threshold $\tau \in [0, 1]$. The resulting routing policy π_τ maps each query to one of the two retrieval modes,

$$\pi_\tau(q) = \begin{cases} M, & s_\phi(q) \geq \tau, \\ N, & s_\phi(q) < \tau, \end{cases} \quad (3.7)$$

where M denotes *Mix* and N denotes *Naive*. Routing is treated as a binary classification problem with `mix_required` as the positive class and

`naive_enough` as the negative class, and the metrics reported below are the standard binary-classification metrics for that labeling. At a chosen threshold these are routing-label accuracy, balanced routing-label accuracy, precision P and recall R on the positive (`mix_required`) class, and the F1 score

$$F_1 = \frac{2PR}{P + R}. \quad (3.8)$$

Throughout this thesis the word accuracy is used in the routing-label sense. It measures whether the policy selected the same mode as the judge-derived binary label. Balanced routing-label accuracy is included because the class distribution is imbalanced and plain routing-label accuracy can be high even for a trivial always-*Naive* policy.

Where bracketed confidence intervals are reported, they are computed using a 1 000-iteration bootstrap over the held-out test set. Each bootstrap sample resamples the test rows with replacement and recomputes the corresponding metric, and reported intervals are the 2.5th and 97.5th percentiles of the resulting bootstrap distribution.

ModernBERT is trained with three random seeds, and the reported ModernBERT router is a seed-averaged ensemble of the three runs, formed by averaging the per-query *Mix* scores. The three individual runs vary little on the aggregate routing-label metrics (standard deviations below 0.02 across seeds), so the ensemble does not mask instability across seeds. The ensemble is then bootstrapped on the test set exactly as the **TF-IDF** baseline is, so both routers carry comparable 95 % bootstrap confidence intervals rather than different dispersion measures.

3.5.2 Offline routing policy

The offline policy evaluation reuses the precomputed retrieval outputs and execution times from Section 3.2.1. For each held-out query, both r_i^N and r_i^M are already available, along with their execution times t_i^N and t_i^M . A routing policy π_τ at threshold τ is then simulated row by row. For each query the policy selects either the *Naive* or the *Mix* precomputed result, and aggregate statistics are computed across the entire held-out split.

Let m denote the number of held-out queries, and let $\mathbf{1}\{\cdot\}$ denote the indicator function that returns 1 when its condition holds and 0 otherwise. Let $r_M(\pi_\tau) = \frac{1}{m} \sum_{i=1}^m \mathbf{1}\{\pi_\tau(q_i) = M\}$ denote the fraction of test queries routed to *Mix* under policy π_τ . Two cost metrics are reported. The mean routed execution time $C(\pi_\tau)$ averages the per-query execution time of whichever

retrieval mode the policy selected,

$$C(\pi_\tau) = \frac{1}{m} \sum_{i=1}^m [\mathbf{1}\{\pi_\tau(q_i) = N\} t_i^N + \mathbf{1}\{\pi_\tau(q_i) = M\} t_i^M]. \quad (3.9)$$

The mean routed model-token cost $U(\pi_\tau)$ is computed in the same way as the mean routed time in Equation 3.9, but with the per-mode token means $\bar{u}^N$ and $\bar{u}^M$ from the instrumented subset (Section 3.2.1) in place of the per-query execution times. The mean routed token cost is therefore an extrapolation from the instrumented subset and is reported alongside the per-query execution time in Chapter 4.

The same simulation is run for two static baseline policies, always-*Naive* and always-*Mix*, which bracket what a static retrieval strategy can achieve on the test set without any per-query decision. They serve as the reference points against which the learned router’s routing-label and routed-cost trade-off is reported.

In addition to mean routed cost, the offline policy evaluation reports two routing error rates with different practical meanings. The under-routing rate, which is the false-negative rate of the binary classifier, is the fraction of test queries where the true label is `mix_required` but the policy routes to *Naive*,

$$\text{under-routing}(\pi_\tau) = \frac{1}{m} \sum_{i=1}^m \mathbf{1}\{y_i = 1, \pi_\tau(q_i) = N\}, \quad (3.10)$$

and the over-routing rate, the false-positive rate, is the fraction of test queries where the true label is `naive_enough` but the policy routes to *Mix*,

$$\text{over-routing}(\pi_\tau) = \frac{1}{m} \sum_{i=1}^m \mathbf{1}\{y_i = 0, \pi_\tau(q_i) = M\}. \quad (3.11)$$

The two error types carry different practical weight. Under-routing is the error type most directly associated with answer-quality risk, because the policy selected *Naive* on a query where *Mix* was judged necessary. Over-routing is primarily a cost error, because the policy selected *Mix* on a query where *Naive* was already judged sufficient. Reporting them separately matters because the headline routing-label metrics count both as errors of the same kind.

Threshold selection uses the validation split only. The threshold reported on test is the one that maximizes F1 on the `mix_required` class on the validation split, and the same threshold is then applied unchanged to test. This prevents the threshold from being tuned on the data it is evaluated on.

Chapter 4

Implementation

This chapter describes the implementation and the experiments. Section 4.1 describes how the two router models were trained. Section 4.2 presents the held-out classification experiment, which evaluates how well each router predicts the judge-derived routing label. Section 4.3 presents the offline routed-cost experiment, which evaluates each router as a thresholded routing policy against the static always-*Naive* and always-*Mix* baselines. Section 4.4 reports the results of both experiments. Section 4.5 discusses what the results imply about the predictability of retrieval-mode selection, where the router fails, and what the practical routing-label and routed-cost trade-off looks like.

4.1 Model training

Both routers are trained on the 900-example training split of the binary routing dataset described in Section 3.3. The validation split is used only for early stopping (ModernBERT) and threshold selection (both routers). The test split is used only for the metrics reported in Section 4.4.

4.1.1 TF-IDF logistic-regression baseline

The baseline is built with scikit-learn [21]. Questions are lowercased and stripped of standard English stop words, then converted into sparse term-frequency vectors with unigram and bigram features (scikit-learn `ngram_range` set to `(1, 2)`), a minimum document frequency of 2, and a vocabulary cap of 20,000 features. The term frequencies are weighted by inverse document frequency [17] learned on the training split. A logistic regression with L_2 regularization (regularization strength $C = 1.0$) is fitted

on top of these features, with class weights set to `balanced` so that the loss on each example is inversely proportional to the frequency of its class in the training data. Optimization uses scikit-learn’s L-BFGS solver [22] with a maximum of 1,000 iterations. Fitting on the 900 training questions takes a few seconds on a single CPU. The same fitted pipeline (vectorizer plus regression head) is then used to score the validation and test splits, producing a probability that each query is *Mix*-beneficial under the judged setup.

4.1.2 ModernBERT router

The ModernBERT router uses the `answerdotai/ModernBERT-base` checkpoint from Hugging Face [16] with a binary classification head. Questions are tokenized with the standard ModernBERT tokenizer and truncated or padded to a maximum length of 128 tokens, which accommodates every question in the 2WikiMultihopQA splits.

Training uses the AdamW optimizer [19] with a learning rate of 2×10^{-5} , a weight decay of 0.01, gradient accumulation of two steps with a per-device batch size of 8 (effective batch size 16), a linear warm-up over the first 10 % of steps followed by linear decay, and gradient clipping at 1.0. The loss is class-balanced cross-entropy with weights α_c inversely proportional to class frequency in the training split, so that the minority `mix_required` class is not down-weighted by the imbalanced training distribution. The model is trained for up to 8 epochs with early stopping based on validation F1 on the `mix_required` class, requiring at least 3 epochs before stopping and using a patience of 3 epochs.

To separate variance from random initialization from variance from finite test data, the ModernBERT router is trained with three different random seeds (7, 13, 42). Each seed independently selects its own best checkpoint on the validation split and its own decision threshold on the validation split. The three runs are combined into a single seed-averaged ensemble (averaging the per-query *Mix* scores), which is the ModernBERT router reported in Section 4.4; its decision threshold is selected on the averaged validation scores. All ModernBERT training and inference was performed on CPU. The model and dataset are small enough that the full fine-tuning run completes in a few hours on a single CPU, which is part of why a lightweight encoder router is attractive as a routing layer.

4.2 Held-out classification experiment

The held-out classification experiment evaluates each router as a binary classifier of judge-derived routing labels on the test split. The router scores are first evaluated threshold-independently, by computing **AUPRC** on the `mix_required` class and **AUROC**. **AUPRC** is the primary metric because it is sensitive to performance on the minority class, which is the class that matters for the routing decision [20].

For thresholded metrics, the decision threshold is selected on the validation split. The selection rule is the threshold τ^* that maximizes F1 on `mix_required` on the validation split, swept over a fixed grid in $[0.05, 0.95]$. The selected τ^* is then applied unchanged to the test split, where routing-label accuracy, balanced routing-label accuracy, precision, recall, and F1 on `mix_required` are reported. A trivial always-*Naive* predictor is included as a simple baseline because the class imbalance makes plain routing-label accuracy easy to inflate.

A non-deployable diagnostic baseline is also included to test whether the learned query-only router improves over a heuristic that routes by a coarse question-category label. The category label is taken directly from the benchmark’s per-question type annotation; it is never visible to the learned routers, which see only the query text. The rule predicts *Mix* for `compositional` questions and *Naive* for all other question types, because `compositional` is the only question type whose *Mix*-beneficial rate exceeds 50% in the training split. This baseline is not a realistic deployment policy because such a category label is unavailable for arbitrary user queries; it is included only as a reference for how much of the routing signal a coarse categorical rule can already deliver, and as a way to measure what the learned query-only routers add on top of it.

For TF-IDF, bracketed values are 95% bootstrap intervals [23] over 1 000 test-set resamples. ModernBERT denotes the seed-averaged ensemble of the three training runs and is bootstrapped identically to TF-IDF.

4.3 Offline routed-cost experiment

The offline routed-cost experiment evaluates each router as a thresholded routing policy along two cost dimensions, following the cost–quality routing-evaluation framing introduced by RouterBench [5]. The precomputed retrieval outputs (the *Naive* and *Mix* answers from Equation 3.1) and per-query

execution times for both LightRAG modes are reused, so that the simulation does not have to re-run any retrieval. A routing policy π_τ at threshold τ sends each test query either to the *Naive* or to the *Mix* precomputed result. The mean routed execution time is computed as in Equation 3.9. The mean routed model-token cost is computed the same way but with per-mode mean token usage from the instrumented subset described in Section 3.2.1 in place of per-query times, giving an extrapolated mean token cost. Under-routing and over-routing rates are computed for each routing policy following Equations 3.10 and 3.11.

For each router, the threshold curve is computed over τ values stepping from 0 to 1 in increments of 0.01, producing a trace of mean routed time and F1 on `mix_required` at each threshold. The same simulation is run for the two static baseline policies, *always-Naive* and *always-Mix*. These two points bracket what a static policy can achieve on the test set and serve as the reference points the learned routers are compared against.

4.4 Results

Tables 4.1 and 4.2 summarize the test-set results: Table 4.1 reports routing-label classification performance, and Table 4.2 reports offline routed cost and policy diagnostics. The asymmetry between the two error types is discussed in Section 4.5.

ModernBERT achieves higher point estimates than **TF-IDF** on every aggregate routing-label metric: **AUPRC**, balanced routing-label accuracy, F1 on the *Mix*-beneficial class, and routing-label accuracy. The bootstrap intervals of the two routers overlap on every threshold-independent metric (Table 4.1), so these gaps are not statistically separated on this 193-row test set; the routing-label-accuracy gap (0.865 vs. 0.788) is the widest and comes closest to separation. The gain comes mainly from higher precision. ModernBERT sends fewer *Naive*-sufficient queries to *Mix*, so its over-routing rate is lower (0.052 vs. 0.155 for **TF-IDF**). **TF-IDF** pulls in the other direction: it has higher *Mix* recall and a lower under-routing rate (0.057 vs. 0.083 for ModernBERT), so it misses fewer *Mix*-beneficial queries, at the cost of routing more unnecessary queries to *Mix*. The two routers therefore sit at opposite ends of the precision-recall trade-off: ModernBERT is more selective and lower-cost; **TF-IDF** is more conservative and spends more tokens to keep recall on the minority class up.

The type-only diagnostic heuristic is strong, which shows that a coarse question-category rule already captures a substantial part of the routing signal.

Its role is diagnostic only, because this heuristic needs a category label that is unavailable for arbitrary user queries. ModernBERT achieves higher point estimates than the categorical heuristic on F1 (0.717 vs. 0.667), precision (0.767 vs. 0.585), and routing-label accuracy (0.865 vs. 0.803), and it does so from query text alone. **TF-IDF** matches the heuristic’s *Mix* recall (0.776) but at lower precision and lower overall accuracy. The cost dimension follows the precision pattern: **TF-IDF** routes at 8,531 mean tokens per query, ModernBERT at 6,666, and always-*Mix* at 17,857, which translates into token savings of 52 % and 63 % relative to the expensive baseline.

Table 4.1: Routing-label classification performance on the held-out test split. All accuracy values are routing-label accuracies (agreement with the judge-derived label); *Bal. acc.* is their class-balanced version and P_{mix} , R_{mix} , $F1_{mix}$ are on the `mix_required` class at the validation-selected threshold. Bracketed values are 95 % bootstrap intervals over 1 000 test-set resamples; ModernBERT is the seed-averaged ensemble of the three runs, bootstrapped identically to TF-IDF. The constant-score baselines have **AUROC** = 0.500 and **AUPRC** equal to the base rate 0.254; the type-only heuristic is non-deployable as it uses the benchmark’s question-category label.

Policy	AUPRC	AUROC	Bal. acc.	P_{mix}	R_{mix}	$F1_{mix}$	Acc.
Always- <i>Naive</i>	0.254	0.500	0.500	0.000	0.000	0.000	0.746
Always- <i>Mix</i>	0.254	0.500	0.500	0.254	1.000	0.405	0.254
Type-only diagnostic	—	—	0.794	0.585	0.776	0.667	0.803
TF-IDF	0.742 [0.610, 0.853]	0.892 [0.836, 0.936]	0.784	0.559	0.776	0.650 [0.547, 0.740]	0.788 [0.731, 0.845]
ModernBERT	0.769 [0.645, 0.889]	0.910 [0.865, 0.949]	0.802 [0.730, 0.872]	0.767	0.673	0.717 [0.607, 0.810]	0.865 [0.819, 0.912]

Table 4.2: Offline routed-cost and policy diagnostics on the held-out test split. For each policy the table gives the fraction of queries routed to *Mix*, the resulting mean per-query time and estimated mean tokens (a mode-weighted average of the per-mode means from the instrumented $N = 15$ subset), the token saving relative to always-*Mix*, and the absolute false-negative (*Under-rout.*) and false-positive (*Over-rout.*) routing-error fractions over all test queries. Under-routing carries the answer-quality risk; over-routing is mainly a cost error. ModernBERT is the seed-averaged ensemble.

Policy	Route-to- <i>Mix</i> fr.	Mean time (s)	Mean tokens	Token saving	Under-rout.	Over-rout.
Always- <i>Naive</i>	0.000	10.95	3 458	80.6 %	0.254	0.000
Always- <i>Mix</i>	1.000	24.42	17 857	0.0 %	0.000	0.746
Type-only diagnostic	0.337	15.59	8 308	53.5 %	0.057	0.140
TF-IDF	0.352	16.98	8 531	52.2 %	0.057	0.155
ModernBERT	0.223	14.11	6 666	62.7 %	0.083	0.052

Figure 4.1 shows the test-set precision–recall curves on the `mix_required` class for the **TF-IDF** baseline and the ModernBERT ensemble. The base rate (the test-set fraction of `mix_required`) is shown as a horizontal dashed

line. The ModernBERT curve sits above the **TF-IDF** curve over most of the recall range, which matches the higher **AUPRC** point estimate in Table 4.1, though the two **AUPRC** bootstrap intervals overlap. The **AUROC** curves in Figure 4.2 are closer together because **AUROC** is less sensitive to performance on the minority class.




Figure 4.1: Precision–recall curves on the held-out test split for the **TF-IDF** baseline and the ModernBERT seed-averaged ensemble. The dashed horizontal line is the test-set base rate of `mix_required`.

Figure 4.3 shows how F1 on `mix_required` varies with the decision threshold on the test split, with the validation-selected threshold of each model marked as a dot. The ModernBERT ensemble selects a validation threshold of $\tau^* = 0.44$, and its test F1 is flat over a wide band of thresholds around this value, which indicates that the F1 plateau is broad and that the choice of operating threshold is not a brittle hyperparameter for these routers.

Figure 4.4 shows the mean routed time on test as the threshold sweeps from 0 (every query routed to *Mix*) to 1 (every query routed to *Naive*). The always-*Naive* and always-*Mix* levels are shown as horizontal dashed lines and act as a sanity check. At $\tau = 0$ every router collapses to always-*Mix*, and at $\tau = 1$ every router collapses to always-*Naive*, which is what the figure shows.

Figure 4.5 combines the two preceding views into a frontier of routing-label F1 against mean routed execution time. Each router traces an operating




Figure 4.2: Receiver-operating-characteristic curves on the held-out test split.




Figure 4.3: F1 on `mix_required` as a function of the decision threshold on the held-out test split. Dots mark the validation-selected threshold for each model: TF-IDF $\tau^* = 0.48$ and the ModernBERT ensemble $\tau^* = 0.44$.

curve in (mean routed time, $F1_{\text{mix}}$) space as τ sweeps. The static *always-Naive* and *always-Mix* policies appear as the two extreme operating points.




Figure 4.4: Mean routed execution time on the held-out test split as the decision threshold sweeps. The two horizontal dashed lines are the *always-Naive* and *always-Mix* reference levels.

Both learned routers achieve higher routing-label F1 than either static policy at substantially lower mean routed time than *always-Mix*. The ModernBERT operating point reaches a higher $F1_{mix}$ than **TF-IDF** at a lower mean routed time.

Table 4.3 reports a per-question-type breakdown on the test split. The four 2WikiMultihopQA question types carry very different fractions of *Mix*-beneficial queries on this split: 2 out of 90 for *comparison*, 8 out of 33 for *bridge_comparison*, 38 out of 65 for *compositional*, and 1 out of 5 for *inference*. The *compositional* slice is the one in which the routing decision matters most. The *inference* slice has too few rows to support stable estimates, so its per-class metrics are reported for completeness only. On *compositional*, ModernBERT achieves higher point estimates than **TF-IDF** on both **AUPRC** (0.834 vs. 0.814) and F1 on *mix_required* (0.810 vs. 0.753). On *bridge_comparison* the two routers trade strengths: **TF-IDF** routes more queries to *Mix* on this slice and reaches higher recall (0.375 vs. 0.125), slightly higher **AUPRC** (0.418 vs. 0.396), and higher F1 (0.316 vs. 0.200), while ModernBERT keeps higher precision (0.500 vs. 0.273) and higher overall slice accuracy (0.758 vs. 0.606).

Figure 4.6 visualizes the **AUPRC** comparison on the two non-degenerate slices.




Figure 4.5: Routing-label F1 on the `mix_required` class versus mean routed execution time on the held-out test split. Each curve traces the operating points obtained by sweeping the decision threshold. Stars mark the validation-selected thresholds. Static *always-Naive* and *always-Mix* policies are shown as reference points.

Table 4.3: Per-question-type breakdown on the held-out test split. Metrics are as in Table 4.1 but computed only on each slice’s rows (threshold metrics at each router’s validation-selected threshold); n and n_{mix} are the slice size and its number of `mix_required` queries. The `comparison` and `inference` slices have too few positives (2 and 1) for stable **AUPRC**/F1 and are shown for completeness only. ModernBERT is the seed-averaged ensemble.

Slice	n	n_{mix}	Model	AUPRC	P_{mix}	R_{mix}	$F1_{mix}$	Acc.
comparison	90	2	TF-IDF	0.277	0.000	0.000	0.000	0.978
			ModernBERT	0.140	0.000	0.000	0.000	0.978
compositional	65	38	TF-IDF	0.814	0.636	0.921	0.753	0.646
			ModernBERT	0.834	0.780	0.842	0.810	0.769
bridge_comparison	33	8	TF-IDF	0.418	0.273	0.375	0.316	0.606
			ModernBERT	0.396	0.500	0.125	0.200	0.758
inference	5	1	TF-IDF	0.333	0.000	0.000	0.000	0.400
			ModernBERT	0.500	0.000	0.000	0.000	0.800

Beyond routing-label accuracy, it is useful to ask how much answer quality the cost saving costs. The judge labels imply, for each policy, how often the routed mode returns a correct answer. Under the assumption that *Mix*




Figure 4.6: **AUPRC** on the `mix_required` class, broken down by 2WikiMultihopQA question type on the held-out test split. ModernBERT bars are the seed-averaged ensemble. The `comparison` and `inference` slices have too few `mix_required` test rows to give an informative **AUPRC** and are omitted.

answers correctly on every query where *Naive* already does — so that the only answer-quality regressions a policy can cause are *Mix*-beneficial queries routed to *Naive* — always-*Naive* answers 74.6 % of the test queries correctly and always-*Mix* answers all of them correctly by construction. The **TF-IDF** router retains 94.3 % and the ModernBERT ensemble 91.7 % of always-*Mix*’s correct answers, at 52 % and 63 % lower estimated token cost respectively. The ModernBERT ensemble thus trades roughly eight percentage points of answer recovery for a large cost reduction, and its answer-quality risk is concentrated entirely in its under-routing rate (0.083). This estimate relies on the stated assumption about *Mix* correctness on `naive_enough` queries; re-judging the *Mix* answers on those queries would remove the assumption and is left to future work.

4.5 Discussion

The retrieval-mode distinction between *Naive* and *Mix* is partly predictable from the query text alone. Both learned routers outperform the always-*Naive* baseline on the minority class and reach operating points in Figure 4.5 that

no static policy at comparable cost can match. The ModernBERT router achieves higher point estimates than **TF-IDF** on **AUPRC**, balanced routing-label accuracy, F1 on the minority class, and routing-label accuracy, although the bootstrap intervals overlap on all of these except routing-label accuracy.

The two routing error types have different practical weight. A false negative routes a *Mix*-beneficial query to *Naive* and so risks an answer-quality regression on a query where *Mix* was judged necessary. A false positive routes a *Naive*-sufficient query to *Mix* and only wastes compute, because *Naive* was already judged sufficient. This asymmetry explains why always-*Mix* scores poorly on routing-label accuracy without necessarily producing worse final answers, and it also explains the gap between **TF-IDF** and ModernBERT in Table 4.2: **TF-IDF** has higher recall and lower under-routing, so it misses fewer *Mix*-beneficial queries; ModernBERT has higher precision and lower over-routing, so it issues fewer unnecessary *Mix* calls.

Routing-label performance and routed cost must therefore be interpreted together. The selected ModernBERT operating point behaves selectively, saving more compute at the cost of accepting somewhat more under-routing risk. The **TF-IDF** operating point behaves conservatively in the opposite direction. Neither extreme is universally correct: which is preferable depends on the relative cost of an unnecessary *Mix* call against an answer-quality regression in a given deployment.

The `compositional` slice is the most informative one because the class distribution is balanced and *Mix* is judged beneficial on most of its queries. On this slice ModernBERT achieves an **AUPRC** of 0.834 against 0.814 for **TF-IDF** and lifts F1 on `mix_required` from 0.753 to 0.810. The gain on `compositional` despite the small training set suggests that the encoder may capture signal beyond the lexical patterns available to the **TF-IDF** baseline. The `bridge_comparison` slice is harder for both routers and is the natural target for follow-up work; on this slice **TF-IDF** reaches a slightly higher **AUPRC** (0.418) and F1 (0.316) by routing more queries to *Mix*, while ModernBERT keeps higher precision and higher slice accuracy at the cost of missing more *Mix*-beneficial queries.

The type-only diagnostic baseline shows that much of the routing signal is already carried by a coarse categorical label: predicting *Mix* for `compositional` questions and *Naive* otherwise already reaches an F1 of 0.667 (Table 4.1), because `compositional` questions dominate the *Mix*-beneficial class on this benchmark. ModernBERT, which sees only the query text and never this benchmark-metadata label, achieves higher point estimates than the categorical rule on F1, precision, and routing-label

accuracy. The router should therefore be read as predicting an operational label from phrasing: what it adds over the type-only reference reflects signal the encoder extracts from how questions are worded, and what it does not exceed reflects benchmark structure rather than a deeper property of the queries.

Stepping back from the routing comparison, the labeling stage is itself informative about the retrieval modes. Of the 2,000 judged queries, 714 (about 36 %) fell in the `none_enough` class, where neither *Naive* nor *Mix* produced a correct answer. These queries are excluded from the routing dataset because routing between two modes cannot help when both fail, but the size of the class is notable in its own right: on a large share of these multi-hop questions the limiting factor is the strength of the retriever, not the choice between its two modes. A stronger retrieval backbone, rather than a better router over *Naive* and *Mix*, is what those queries would need.

The benchmark used here is multi-hop by construction, so every query is a comparatively hard case for vector-only retrieval and the `mix_required` rate is correspondingly high. In a deployment such as Ericsson’s, where a substantial fraction of queries are single-hop factual look-ups that *Naive* retrieval already handles, more traffic would fall in the `naive_enough` class. An accurate router would then send an even larger share of queries to the low-cost path, so the cost saving reported here is, if anything, a conservative estimate of what query-only routing could deliver on a less uniformly multi-hop workload.

The per-question-type breakdown and the type-only diagnostic baseline both rely on the benchmark’s per-question category label. In a real deployment that label does not exist; users send free-form queries and the routing layer sees only the query text. The category-based slices are useful as a diagnostic device because they show which subpopulations the router handles well and which it does not, and they help separate signal that comes from coarse structural cues (such as whether the question is a compositional multi-hop lookup) from signal that comes from finer phrasing. They are not a deployment recipe; the learned routers approximate the same structure from the query text alone, which is the only input available outside the benchmark.

The same framing helps interpret the contrast between the `compositional` and `bridge_comparison` slices. The fact that ModernBERT is strong on `compositional` and weaker on `bridge_comparison` indicates which phrasings the encoder has learned to recognize on this benchmark, and does not extend to a claim that `compositional` queries are easier to route in production. Within this benchmark, the `compositional` slice contains more rows and the rows share a more consistent surface form.

Always-Mix illustrates why the two result tables must be read together. It is label-incorrect on *Naive-sufficient* rows because it uses more retrieval than necessary, although the underlying *Mix* answer on those rows can still be correct. The absence of a direct final-answer-accuracy comparison for each routed policy is treated as a limitation in Section 5.2.3.

The reported metrics are still informative for adaptive routing. Routing-label F1 measures whether the router selects the minimum judged sufficient mode. Mean routed time and estimated token usage measure the cost of doing so. The under-routing and over-routing rates separate answer-risk errors from cost-waste errors. Together they describe how often the router agrees with the judge and how expensive that agreement is.

The held-out test set has 193 queries. This is large enough to compute meaningful bootstrap intervals and too small to resolve small differences between the routers with confidence. When the ModernBERT ensemble is bootstrapped on the same test set as **TF-IDF**, the two routers' 95 % intervals overlap on every threshold-independent metric: the **AUPRC** intervals ($[0.610, 0.853]$ for **TF-IDF** and $[0.645, 0.889]$ for ModernBERT) overlap heavily, and the same holds for **AUROC** and balanced routing-label accuracy. The routing-label-accuracy intervals ($[0.731, 0.845]$ for **TF-IDF** and $[0.819, 0.912]$ for ModernBERT) come closest to separation but still overlap. The point estimates therefore favour ModernBERT consistently, yet none of the individual gaps is statistically significant at this sample size. A larger benchmark would be needed before stronger claims can be made.

Finally, the cost picture is sharper in model tokens than in execution time. The two static baselines bracket a mean routed range of 10.95 to 24.42 seconds per query in execution time, and 3,458 to 17,857 tokens per query in model-token usage. The learned routers select operating points around 14 to 17 seconds per query (a 30–42 % reduction relative to *always-Mix*) and 6,700 to 8,500 tokens per query (a 52–63 % reduction relative to *always-Mix*), at non-trivial `mix_required` F1. The token reduction is larger than the time reduction because *Mix* adds one extra LLM call per query and a substantially larger assembled context [3, 12], while execution time is partly absorbed by network latency that does not scale with token volume. The routing decision itself contributes zero LLM tokens because neither router calls a generative model, and its inference cost is expected to be negligible relative to retrieval cost on either dimension.

Chapter 5

Conclusions and future work

This chapter summarizes the conclusions from the experiments (Section 5.1), discusses the main limitations and natural avenues for future work (Section 5.2), and reflects on the broader context of the project (Section 5.3).

5.1 Conclusions

This thesis studied whether the choice between LightRAG’s low-cost *Naive* retrieval mode and its more expensive *Mix* retrieval mode can be predicted from the query text alone on the 2WikiMultihopQA benchmark using a lightweight encoder-only classifier instead of a large language-model controller. Routing labels were derived by running both retrieval modes on each benchmark question and asking an **LLM-as-a-Judge** [18] to compare each mode’s answer against the gold answer. The label `mix_required` is operational: it captures whether, under the chosen LightRAG configuration and judge prompt, *Mix* was judged correct when *Naive* was not, and it does not isolate which component of *Mix* drove the improvement.

The research question asked, within a fixed Hybrid **RAG** setting, how well a lightweight query-only classifier predicts a judge-derived routing label for the minimum sufficient retrieval mode, and how using that classifier as an offline routing policy affects routed cost relative to static *Naive* and *Mix* baselines. The remainder of this section answers the two sub-questions in turn.

The first sub-question concerns classifier accuracy on the routing label. On the held-out test split of 193 queries, the ModernBERT ensemble achieves higher point estimates than the **TF-IDF** baseline on every aggregate routing-label metric: **AUPRC** (0.769 vs. 0.742), balanced routing-label accuracy

(0.802 vs. 0.784), F1 on the minority `mix_required` class (0.717 vs. 0.650), and routing-label accuracy (0.865 vs. 0.788). When both routers are bootstrapped on the same test set, however, their 95 % intervals overlap on every threshold-independent metric, so none of these gaps is statistically significant at 193 test queries; the routing-label-accuracy gap is the widest and comes closest to separation. The practical answer to the first sub-question is therefore that the fine-tuned encoder predicts the routing label at least as well as the **TF-IDF** baseline, and consistently a little better in point estimate, but the present test set is too small to certify a strict improvement. The two routers place themselves at opposite ends of the precision-recall trade-off on the `mix_required` class: **TF-IDF** has lower under-routing (0.057 vs. 0.083), and ModernBERT has lower over-routing (0.052 vs. 0.155). ModernBERT also beats the non-deployable type-only heuristic — which predicts *Mix* for `compositional` questions and *Naive* otherwise, reaching F1 0.667 and routing-label accuracy 0.803 — on F1, precision, and routing-label accuracy, using only the query text rather than benchmark metadata.

The second sub-question concerns routed cost relative to static baselines. Both learned routers reach test-set operating points around 14 to 17 seconds of mean routed execution time, compared with 10.95 seconds for always-*Naive* and 24.42 seconds for always-*Mix*, and achieve higher routing-label F1 than either static policy at substantially lower mean routed time than always-*Mix*. Measured in model tokens, they route at roughly 6,700 to 8,500 tokens per query, compared with 3,458 for always-*Naive* and 17,857 for always-*Mix*. The selected ModernBERT operating point reduces estimated routed token cost by 63 % relative to always-*Mix*, and the selected **TF-IDF** operating point by 52 %. The answer to the second sub-question is therefore that query-only routing recovers most of the cost of always-*Mix* — up to a 42 % reduction in mean time and a 63 % reduction in estimated tokens for the ModernBERT ensemble — while, under the assumption that *Mix* answers correctly wherever *Naive* does, still producing a correct final answer on about 92 % of the queries that always-*Mix* answers correctly. The routing decision itself adds negligible cost on either dimension because neither router calls a generative model; ModernBERT’s single encoder forward pass costs more than computing the **TF-IDF** features and logistic-regression score, but both are negligible next to the retrieval and generation they gate.

Taken together, the two answers indicate that the *Naive*-sufficient versus *Mix*-beneficial distinction is partly predictable from query text alone in this setting, and that a lightweight encoder classifier extracts at least as much of that predictability as a **TF-IDF** logistic-regression baseline — a little more in

point estimate, though within overlapping confidence intervals — and more than a coarse question-category heuristic. The reported accuracy values are routing-label accuracies, and they are distinct from independently measured final answer accuracies. Within that scope, the result supports learned query routing as a promising cost-control layer for Hybrid **RAG**: the router can often identify when the expensive path is unnecessary, and it therefore reduces routed cost without always defaulting to the most expensive retrieval mode.

5.2 Limitations and future work

5.2.1 Token-cost measurement scope

Model-token cost in this thesis is estimated using per-mode mean token usage from an instrumented subset of $N = 15$ queries spanning all four 2WikiMultihopQA question types. This subset provides a preliminary estimate of per-mode mean token usage, and full per-query logging across the 1,286-row routing dataset would be needed for exact routed token costs and uncertainty estimates. The reported token costs should therefore be read as estimated routed token usage from per-mode means, and not as exact measured token usage for every held-out query.

Per-query token logging would also make it possible to report bootstrap intervals for token cost on the same footing as for execution time, and to study whether token usage varies systematically across question types or routing labels.

5.2.2 Judge labels

Routing labels come from an **LLM-as-a-Judge** and are used as judged supervision, not as absolute ground truth. The judge sees the gold answer and the same prompt is used consistently across the dataset [18], but judge errors may still propagate into the labels. The risk is largest for borderline answers, partially correct answers, and answers that use different surface forms from the gold answer.

The label `mix_required` should be read operationally: it captures cases in which, under the chosen LightRAG configuration and judge prompt, *Mix* was judged correct when *Naive* was not, and it does not isolate which component of *Mix* caused the improvement. *Mix* differs from *Naive* in graph traversal, context assembly, and retrieval breadth, and the label aggregates over all of these.

Follow-up work would include validating a sample of judge decisions manually, comparing multiple judge models, and tightening the answer-normalization rules where possible.

5.2.3 Final answer accuracy of routed policies

The offline policy evaluation reports routing-label performance and routed cost. It does not independently re-judge the final answer selected by each routed policy. This matters because false-positive *Mix* decisions are counted as routing-label errors, although the *Mix* answer for the same query may still be correct. False-negative *Mix* decisions are the errors most directly associated with answer-quality risk, because the router selects *Naive* on a query where *Mix* was judged necessary.

A future evaluation should retain per-mode correctness indicators for every query, or re-run the judge on the answer selected by each policy, so that judged final answer accuracy can be reported alongside routed cost. With per-policy answer correctness available, the main trade-off would become answer correctness against token cost.

5.2.4 Generalizability

The experiments use one Hybrid **RAG** framework (LightRAG), two of its retrieval modes (*Naive* and *Mix*), and one public benchmark (2WikiMulti-hopQA). The held-out test set has 193 rows, which is enough to compute bootstrap intervals and not enough to resolve small differences between routers. The conclusions therefore apply most directly to the studied setting. Applying the same methodology to other Hybrid **RAG** frameworks (for example GraphRAG [2] or hybrid retrieval modes in other systems) and to other multi-hop benchmarks (HotpotQA, MuSiQue [9, 10]) is a natural way to test how robust the routing signal is. Extending the routing labels to industrial corpora (for example Ericsson internal documentation) was considered but is blocked by the absence of per-query gold answers, which is what made the judge-labeling protocol possible on 2WikiMultihopQA in the first place.

The strong type-only diagnostic baseline also indicates that part of the routing signal may be benchmark-specific. The results should not be read as evidence that the same router would transfer unchanged to arbitrary industrial queries without additional validation.

5.2.5 Routing beyond the query

The router used in this thesis is query-only by construction. A router that also has access to some lightweight signal from the corpus or from a low-cost first retrieval pass could in principle make a better decision, at the cost of a slightly more expensive routing layer. A natural follow-up is to compare the query-only router against a two-stage scheme that first runs the *Naïve* mode, inspects whether the retrieved chunks contain the entities referenced in the query, and escalates to *Mix* only when this signal looks weak. This is similar to how Adaptive-RAG conditions on inexpensive intermediate signals at the language-model level [4]. The routing layer would remain inexpensive (one round of vector retrieval instead of an expensive graph traversal) and could materially improve the harder slices, in particular `bridge_comparison`.

5.2.6 Calibration and abstention

The router currently outputs a single thresholded decision. Routing under cost constraints is a natural setting for selective prediction [24] and calibrated probabilities [25], where the router could abstain on low-confidence queries and either escalate to a more expensive but more reliable controller or run both retrieval modes and combine the answers. The calibration of the ModernBERT router’s softmax outputs was not investigated in this work and is a natural follow-up.

5.3 Reflections

This thesis is centered on making a Hybrid **RAG** system more resource-efficient. The most direct connection is to the **United Nations (UN) Sustainable Development Goal (SDG)** number 12 (responsible consumption and production). Graph-enhanced retrieval is significantly more compute-intensive than vector retrieval [3, 12], and a router that avoids applying it to queries that do not need it reduces measured compute-related cost proxies such as routed token usage and execution time. This may reduce associated energy use in deployment, although energy consumption was not measured directly. There is also a softer connection to **SDG** number 9 (industry, innovation, and infrastructure), in that adaptive routing is a building block for industrial **RAG** deployments such as the one motivating this work at Ericsson, where serving company-internal documentation at scale is a real operational concern.

The work has some ethical aspects worth flagging. The routing labels are

produced by an **LLM-as-a-Judge** that has access to the gold answer. This is a methodologically reasonable but imperfect protocol. It is acceptable for the present study for two reasons: the judge decides a constrained, reference-grounded comparison against the gold answer rather than an open-ended quality score, which limits how far it can drift from a correct reading; and the prediction target is deliberately the judge’s operational routing decision, not a human’s notion of answer quality. Validating a sample of judge decisions against human judgments, and measuring how often the two disagree, is left to future work, and the thesis is explicit throughout about not treating these labels as ground truth. Building a system that decides on a per-query basis which retrieval path to use could in principle interact with downstream notions of answer quality in ways that are hard to anticipate. If the router systematically routes some kind of question to the lower-cost path, that kind of question could end up consistently worse-served, even when the average behavior of the system improves. The per-question-type breakdown in Section 4.4 is one way to surface this kind of differential behavior, and any future work that scales this approach to user-facing settings should keep similar per-segment analyses in view.

References

- [1] Y. Gao, Y. Xiong, X. Gao, K. Jia, J. Pan, Y. Bi, Y. Dai, J. Sun, M. Wang, and H. Wang, “Retrieval-augmented generation for large language models: A survey,” 2024. [Online]. Available: <https://arxiv.org/abs/2312.10997>
- [2] D. Edge, H. Trinh, N. Cheng, J. Bradley, A. Chao, A. Mody, S. Truitt, D. Metropolitansky, R. O. Ness, and J. Larson, “From local to global: A graph rag approach to query-focused summarization,” 2025. [Online]. Available: <https://arxiv.org/abs/2404.16130>
- [3] Z. Guo, L. Xia, Y. Yu, T. Ao, and C. Huang, “Lightrag: Simple and fast retrieval-augmented generation,” 2025. [Online]. Available: <https://arxiv.org/abs/2410.05779>
- [4] S. Jeong, J. Baek, S. Cho, S. J. Hwang, and J. Park, “Adaptive-rag: Learning to adapt retrieval-augmented large language models through question complexity,” in *Proceedings of the 2024 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies (Volume 1: Long Papers)*, K. Duh, H. Gomez, and S. Bethard, Eds. Mexico City, Mexico: Association for Computational Linguistics, Jun. 2024. doi: 10.18653/v1/2024.naacl-long.389 pp. 7036–7050. [Online]. Available: <https://aclanthology.org/2024.naacl-long.389/>
- [5] Q. J. Hu, J. Bieker, X. Li, N. Jiang, B. Keigwin, G. Ranganath, K. Keutzer, and S. K. Upadhyay, “Routerbench: A benchmark for multi-llm routing system,” 2024. [Online]. Available: <https://arxiv.org/abs/2403.12031>
- [6] A. Asai, Z. Wu, Y. Wang, A. Sil, and H. Hajishirzi, “Self-rag: Learning to retrieve, generate, and critique through self-reflection,” 2023. [Online]. Available: <https://arxiv.org/abs/2310.11511>

- [7] X. Ho, A.-K. Duong Nguyen, S. Sugawara, and A. Aizawa, “Constructing a multi-hop QA dataset for comprehensive evaluation of reasoning steps,” in *Proceedings of the 28th International Conference on Computational Linguistics*, D. Scott, N. Bel, and C. Zong, Eds. Barcelona, Spain (Online): International Committee on Computational Linguistics, Dec. 2020. doi: 10.18653/v1/2020.coling-main.580 pp. 6609–6625. [Online]. Available: <https://aclanthology.org/2020.coling-main.580/>
- [8] P. Lewis, E. Perez, A. Piktus, F. Petroni, V. Karpukhin, N. Goyal, H. Kuttler, M. Lewis, W.-t. Yih, T. Rocktäschel, S. Riedel, and D. Kiela, “Retrieval-augmented generation for knowledge-intensive NLP tasks,” in *Advances in Neural Information Processing Systems*, vol. 33, 2020, pp. 9459–9474. [Online]. Available: <https://arxiv.org/abs/2005.11401>
- [9] Z. Yang, P. Qi, S. Zhang, Y. Bengio, W. W. Cohen, R. Salakhutdinov, and C. D. Manning, “HotpotQA: A dataset for diverse, explainable multi-hop question answering,” in *Proceedings of the 2018 Conference on Empirical Methods in Natural Language Processing*, 2018. doi: 10.18653/v1/D18-1259 pp. 2369–2380. [Online]. Available: <https://aclanthology.org/D18-1259/>
- [10] H. Trivedi, N. Balasubramanian, T. Khot, and A. Sabharwal, “MuSiQue: Multihop questions via single-hop question composition,” *Transactions of the Association for Computational Linguistics*, vol. 10, pp. 539–554, 2022. doi: 10.1162/tacl_a_00475. [Online]. Available: <https://aclanthology.org/2022.tacl-1.31/>
- [11] V. Karpukhin, B. Oguz, S. Min, P. Lewis, L. Wu, S. Edunov, D. Chen, and W.-t. Yih, “Dense passage retrieval for open-domain question answering,” in *Proceedings of the 2020 Conference on Empirical Methods in Natural Language Processing*, B. Webber, T. Cohn, Y. He, and Y. Liu, Eds. Online: Association for Computational Linguistics, Nov. 2020. doi: 10.18653/v1/2020.emnlp-main.550 pp. 6769–6781. [Online]. Available: <https://aclanthology.org/2020.emnlp-main.550/>
- [12] L. Stamatescu, “GraphRAG costs explained: What you need to know,” Microsoft Foundry Blog, Aug. 2024, accessed 2026-05-04. [Online]. Available: <https://techcommunity.microsoft.com/blog/azure-ai-foundry-blog/graphrag-costs-explained-what-you-need-to-know/4207978>

- [13] A. Vaswani, N. Shazeer, N. Parmar, J. Uszkoreit, L. Jones, A. N. Gomez, Ł. Kaiser, and I. Polosukhin, “Attention is all you need,” in *Advances in Neural Information Processing Systems*, vol. 30, 2017, pp. 5998–6008. [Online]. Available: <https://papers.nips.cc/paper/7181-attention-is-all-you-need>
- [14] J. Devlin, M.-W. Chang, K. Lee, and K. Toutanova, “BERT: Pre-training of deep bidirectional transformers for language understanding,” in *Proceedings of the 2019 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies*, 2019. doi: 10.18653/v1/N19-1423 pp. 4171–4186. [Online]. Available: <https://aclanthology.org/N19-1423/>
- [15] A. Benayas, M. A. Sicilia, and M. Mora-Cantallos, “A comparative analysis of encoder only and decoder only models in intent classification and sentiment analysis: navigating the trade-offs in model size and performance,” *Language Resources and Evaluation*, vol. 59, no. 3, pp. 2007–2030, 2024. doi: 10.1007/s10579-024-09796-y. [Online]. Available: <https://doi.org/10.1007/s10579-024-09796-y>
- [16] B. Warner, A. Chaffin, B. Clavié, O. Weller, O. Hallström, S. Taghadouini, A. Gallagher, R. Biswas, F. Ladhak, T. Aarsen, N. Cooper, G. Adams, J. Howard, and I. Poli, “Smarter, better, faster, longer: A modern bidirectional encoder for fast, memory efficient, and long context finetuning and inference,” 2024. [Online]. Available: <https://arxiv.org/abs/2412.13663>
- [17] G. Salton and C. Buckley, “Term-weighting approaches in automatic text retrieval,” *Information Processing & Management*, vol. 24, no. 5, pp. 513–523, 1988. doi: 10.1016/0306-4573(88)90021-0
- [18] L. Zheng, W.-L. Chiang, Y. Sheng, S. Zhuang, Z. Wu, Y. Zhuang, Z. Lin, Z. Li, D. Li, E. P. Xing, H. Zhang, J. E. Gonzalez, and I. Stoica, “Judging llm-as-a-judge with mt-bench and chatbot arena,” 2023. [Online]. Available: <https://arxiv.org/abs/2306.05685>
- [19] I. Loshchilov and F. Hutter, “Decoupled weight decay regularization,” in *International Conference on Learning Representations*, 2019. [Online]. Available: <https://openreview.net/forum?id=Bkg6RiCqY7>
- [20] T. Saito and M. Rehmsmeier, “The precision-recall plot is more informative than the ROC plot when evaluating binary classifiers

- on imbalanced datasets,” *PLOS ONE*, vol. 10, no. 3, p. e0118432, 2015. doi: 10.1371/journal.pone.0118432. [Online]. Available: <https://doi.org/10.1371/journal.pone.0118432>
- [21] F. Pedregosa, G. Varoquaux, A. Gramfort, V. Michel, B. Thirion, O. Grisel, M. Blondel, P. Prettenhofer, R. Weiss, V. Dubourg, J. Vanderplas, A. Passos, D. Cournapeau, M. Brucher, M. Perrot, and E. Duchesnay, “Scikit-learn: Machine learning in Python,” *Journal of Machine Learning Research*, vol. 12, pp. 2825–2830, 2011.
- [22] D. C. Liu and J. Nocedal, “On the limited memory BFGS method for large scale optimization,” *Mathematical Programming*, vol. 45, no. 1, pp. 503–528, 1989. doi: 10.1007/BF01589116
- [23] B. Efron, “Bootstrap methods: Another look at the jackknife,” *The Annals of Statistics*, vol. 7, no. 1, pp. 1–26, 1979. doi: 10.1214/aos/1176344552
- [24] Y. Geifman and R. El-Yaniv, “Selective classification for deep neural networks,” 2017. [Online]. Available: <https://arxiv.org/abs/1705.08500>
- [25] C. Guo, G. Pleiss, Y. Sun, and K. Q. Weinberger, “On calibration of modern neural networks,” in *Proceedings of the 34th International Conference on Machine Learning*, 2017, pp. 1321–1330. [Online]. Available: <https://proceedings.mlr.press/v70/guo17a.html>

Appendix A

Source code

The implementation code, the data-preparation pipeline, the training scripts, and the evaluation scripts used in this thesis are publicly available at https://github.com/Dantecool02/Thesis_ericsson.

Appendix B

Experimental configuration

This appendix lists the configuration used to produce the routing labels, train the routers, and compute the reported routed-cost estimates. Table B.1 summarizes the LightRAG, data, and judge settings. Table B.2 summarizes the router training settings. Section B.1 maps each pipeline stage onto a script in the public code repository (Appendix A), and Section B.2 reproduces the judge prompt verbatim.

B.1 Pipeline scripts

The eight pipeline stages used in this thesis map one-to-one onto scripts in the public code repository (Appendix A). Table B.3 lists the script for each stage in the order it is executed. Running the scripts in this order reproduces the data, the routers, and the figures used in this thesis from the raw 2WikiMultihopQA download.

B.2 Judge prompt

The judge receives the question, the gold answer, the *Naive* answer, and the *Mix* answer in a JSON payload, and is asked to return exactly one of the three class labels `naive_enough`, `mix_required`, or `none_enough`. The system prompt is reproduced verbatim below; the label semantics that the judge enforces are added below the verbatim prompt as inline annotations and are not part of the prompt itself.

```
You are a strict evaluator for a retrieval-routing
dataset.
```

Table B.1: LightRAG, data, and judge configuration used in this thesis. Each row names one configurable component of the pipeline and gives the exact value used in all reported experiments. *Item* is the name of the component (model checkpoint, framework version, retrieval parameter, or instrumentation choice); *Value* is the literal value passed to the corresponding API or library. The same generator model and embedding model are used during indexing and at query time. All Gemini calls use a temperature of 0.0.

Item	Value
Generator model	gemini-3-flash-preview (Gemini API)
Judge model	gemini-3-flash-preview (Gemini API)
Embedding model	models/gemini-embedding-001 (768-dim)
LightRAG version	1.4.10 (vendored fork in external/-LightRAG)
LightRAG chunk token size	1 200 tokens
LightRAG chunk overlap	100 tokens
<i>Naive</i> retrieval parameters	top_k=3, chunk_top_k=3
<i>Mix</i> retrieval parameters	mode="mix", chunk_top_k=3; entity-relation graph built during indexing
LightRAG tokenizer	gpt-4o-mini via tiktoken
Generator temperature	0.0
Judge temperature	0.0
Token logging method	Gemini API usage_metadata per LLM call, plus tiktoken estimates for embedding calls when the API does not report tokens
Token measurement sub-set	$N = 15$ queries spanning all four 2WikiMultihopQA question types
Judge prompt location	Section B.2

You must output exactly one class label and nothing else

.

Allowed labels:

- naive_enough
- mix_required
- none_enough

Label semantics (annotation, not part of the prompt):

Table B.2: Router training configuration used in this thesis. Each row names one router hyperparameter or training-pipeline setting and gives the exact value used in all reported experiments. *Item* is the name of the hyperparameter or setting; *Value* is the value passed to scikit-learn (for the **TF-IDF** baseline) or Hugging Face Transformers (for the ModernBERT router). All three ModernBERT seeds use the values listed here; only the random seed differs across the three runs.

Item	Value
TF-IDF features	Unigrams and bigrams, lowercase, English stop-words removed, min document frequency 2, vocabulary cap 20 000
TF-IDF classifier	Logistic regression, L_2 regularization with $C = 1.0$, balanced class weights, scikit-learn L-BFGS solver, max 1 000 iterations
ModernBERT checkpoint	answerdotai/ModernBERT-base
ModernBERT max sequence length	128 tokens
ModernBERT learning rate	2×10^{-5}
ModernBERT weight decay	0.01
ModernBERT batch size	Per-device 8, gradient accumulation 2, effective batch 16
ModernBERT optimizer	AdamW with linear warm-up over 10 % of steps and linear decay; gradient clip 1.0
ModernBERT max epochs	8, early stopping on validation $F1_{\text{mix}}$ with patience 3, minimum 3 epochs
ModernBERT seeds	7, 13, 42
Split sizes	900 train / 193 validation / 193 test
Split strategy	Joint stratification by question type and routing label
Hardware	Single CPU; no GPU used for ModernBERT training or inference

```
# naive_enough -> the Naive answer matches the gold
# answer
# mix_required -> the Naive answer does not match the
# gold answer
# but the Mix answer does
# none_enough -> neither answer matches the gold answer;
```

Table B.3: Pipeline stages and the script that implements each stage. *Stage* names a logical step of the data and modelling pipeline (sample preparation, indexing, dual-mode execution, token instrumentation, judging, dataset construction, router training, evaluation, or figure generation). *Script* is the path of the Python file in the public code repository (Appendix A) that implements the stage; all paths are relative to the repository root. Running the scripts in the order shown reproduces the data, routers, and figures used in this thesis from the raw 2WikiMultihopQA download.

Stage	Script
Prepare 2WikiMultihopQA samples	src/dataset_tools/prepare_2wiki_samples.py
Index the corpus in Lightrag	src/data_synthesis/index_2wiki_lightrag.py
Run both retrieval modes on every queryable sample	src/data_synthesis/query_lightrag.py
Instrument token usage on the $N = 15$ subset	src/data_synthesis/query_lightrag_token_usage.py
Generate routing labels with LLM-as-a-Judge	src/data_synthesis/judge_query_results.py
Build the binary routing dataset and joint-stratified splits	src/router/prepare_router_dataset.py
Train the TF-IDF baseline router	src/router/train_baseline_router.py
Train the ModernBERT routers (seeds 7, 13, 42)	src/router/train_modernbert_router.py
Run held-out classification, routed-cost simulation, per-question-type breakdown, and bootstrap CIs for the seed-averaged ensemble	src/router/thesis_analysis_ensemble.py
Generate the figures used in this thesis	src/router/thesis_figures_ensemble.py

```
# these rows are excluded from the binary routing
dataset
```